

# InterProcess Communication

---

Sockets and Port

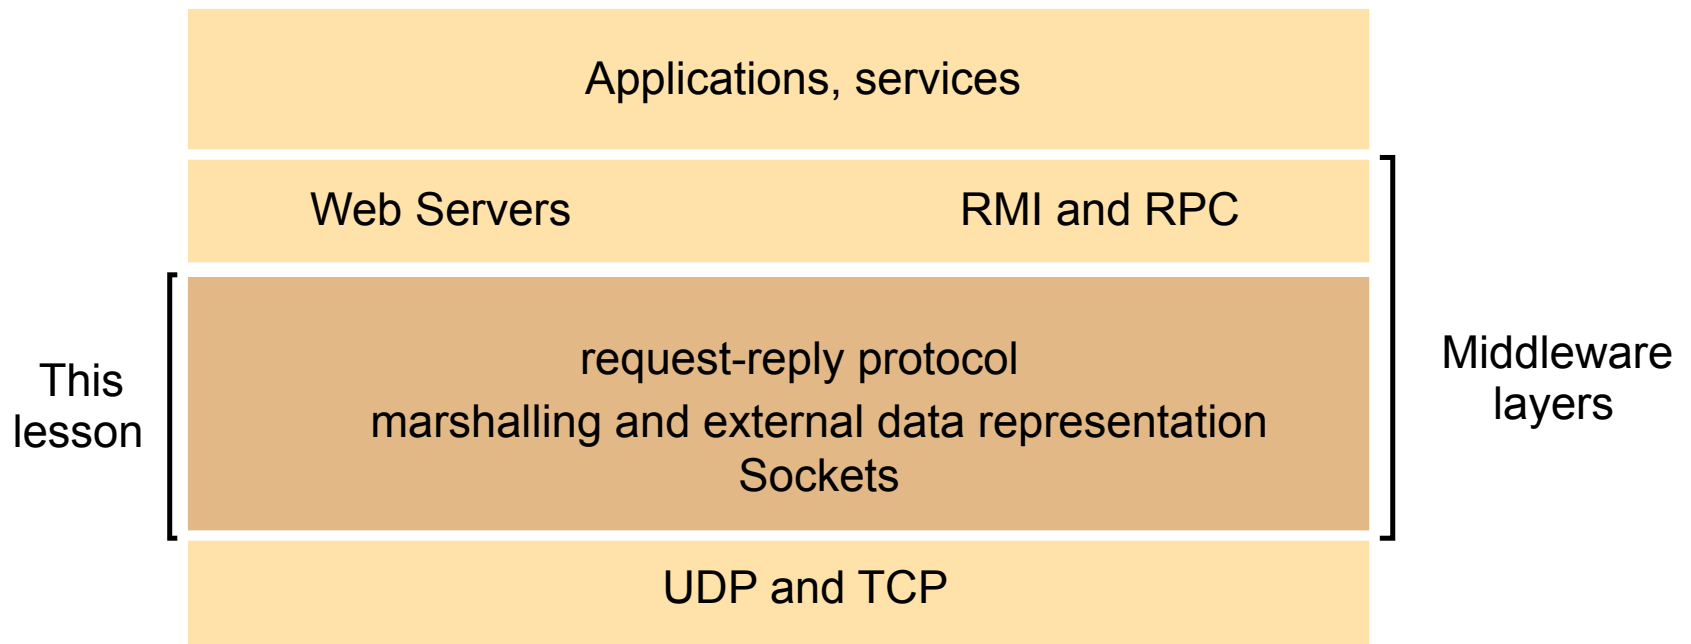
Java Sockets

Marshalling/Unmarshalling data

Request-Reply protocol

HTTP protocol

# Middleware layers



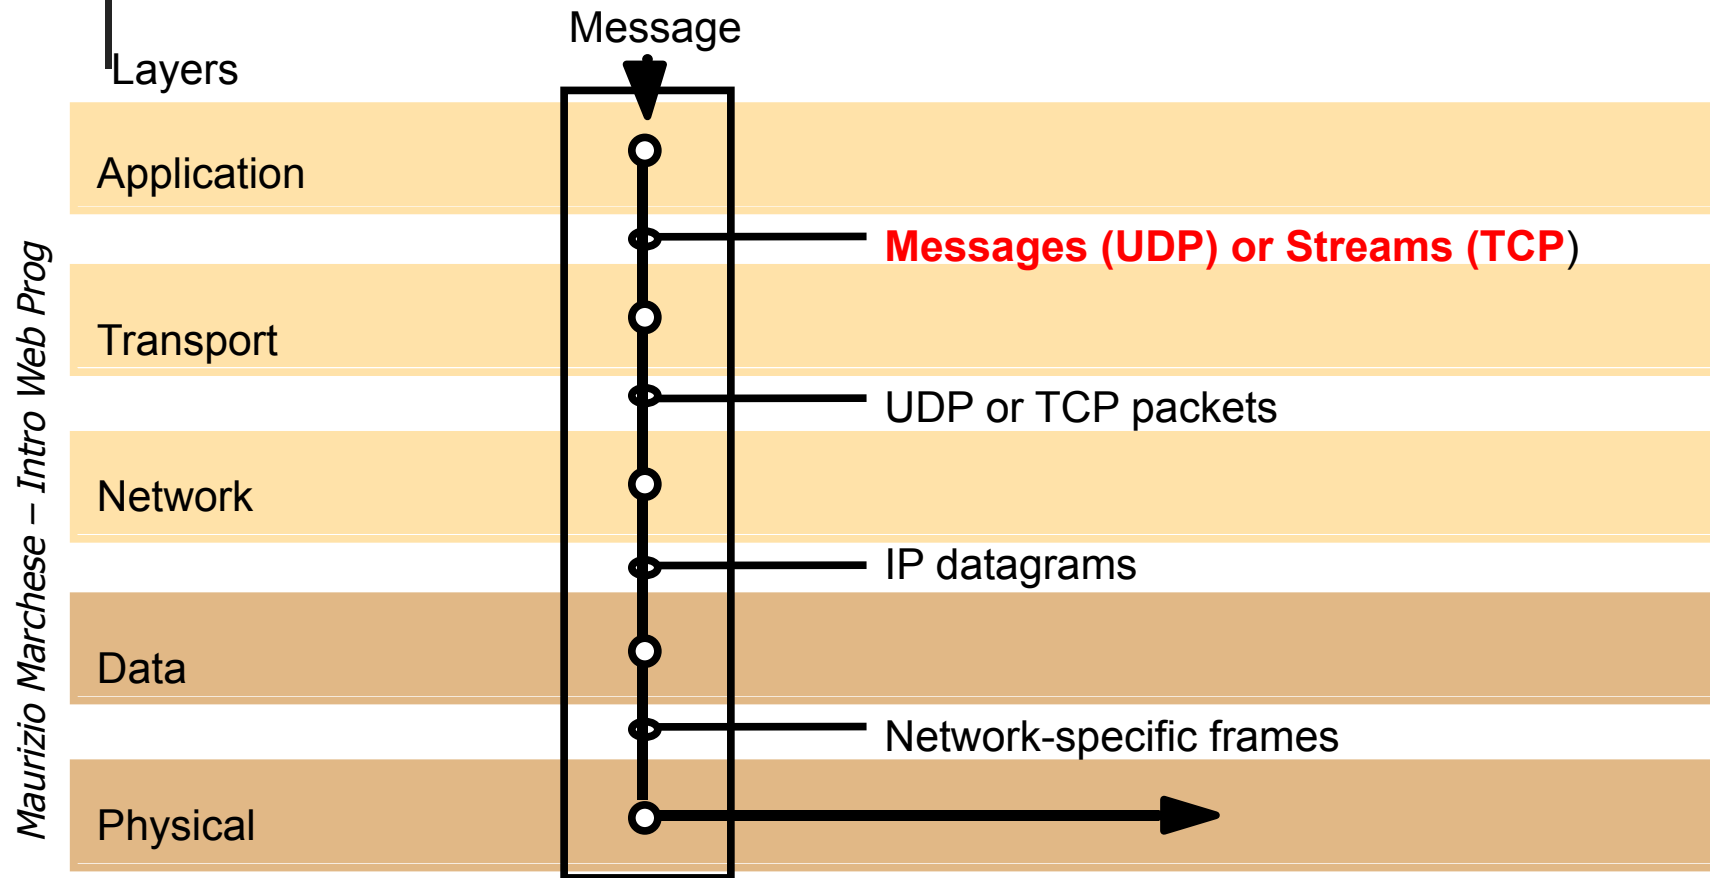


# Clients and Servers


---

- Quando si scrivono **network applications**, è comune parlare di clients e servers. La distinzione è piuttosto vaga, ma sostanzialmente possiamo pensare che :
  - chi inizia la conversazione è il **client**.
  - chi accetta la richiesta di parlare e' il **server**.
- Per i nostri scopi, la più importante differenza tra client e server è che
  - **il client può in qualunque istante creare un canale di comunicazione** di rete (per es. socket) per iniziare una conversazione con una server application,
  - **mentre un server si deve preparare ad ascoltare in anticipo** per possibili conversazioni in arrivo

# TCP/IP layers



# Sockets and ports



Internet address = 138.37.94.248

Internet address = 138.37.88.249



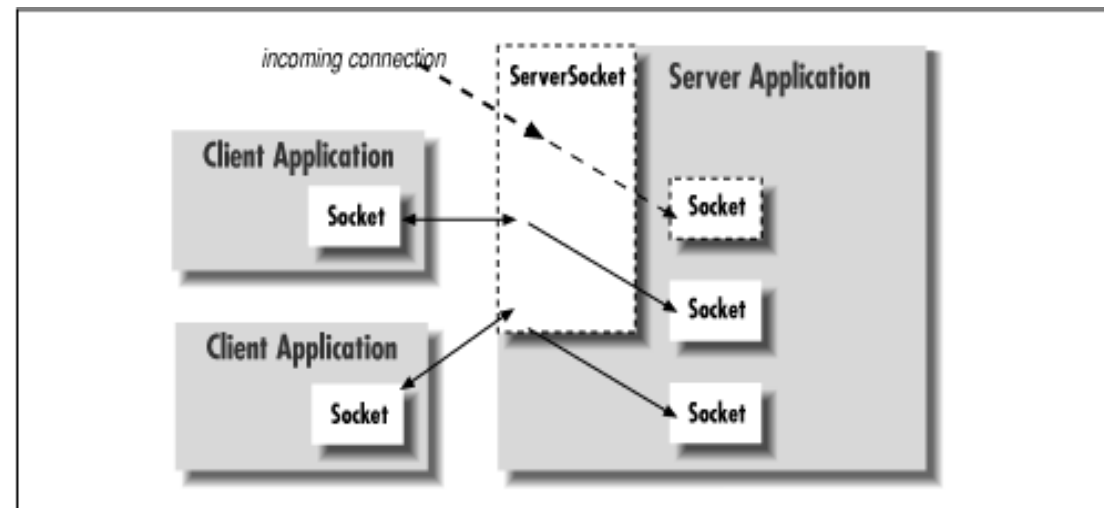
# Java Sockets

---

- Java supporta un accesso semplificato e object-oriented alle sockets.
  - Questo rende la network comunicazione di rete sensibilmente più semplice.
  - Parlare ad un'altra applicazione (locale o remota) è semplice come leggere un file o ottenere un input da un utente.
- La **java.net.Socket** class rappresenta un singolo lato di una connessione socket indifferentemente su un client o su un server.

# Java Sockets

- Un server usa la **java.net.ServerSocket** class per attendere connessioni da clients. Un applicazione agente come server crea un ServerSocket object e attende, bloccato in una chiamata al suo metodo **accept()**, finché non giunge una connessione. A quel punto, il metodo **accept()** crea un Socket object che il server usa per comunicare con il client.
- Un server mantiene molte conversazioni simultaneamente. C'è una sola ServerSocket, mentre c'è un oggetto Socket attivo per ogni cliente.





# Server port

---

- Una applicazione server ascolta su una porta predefinita mentre attende una connessione

```
ServerSocket listenSocket =  
    new ServerSocket(7896)
```

- Un client ha bisogno di due informazioni per connettersi a un server su Internet:
  - un nome di host (per recuperare l'indirizzo del server)
  - un numero di porta (per individuare il processo sulla macchina server).

```
int serverPort = 7896;  
s = new Socket("127.0.0.1", serverPort );
```

- I clients selezionano il numero di porta corrispondente al servizio desiderato.
  - I numeri di porta sono codificati nelle RFC (Es. Telnet 23, FTP 21, ecc.), ma possono anche essere scelti (quasi) arbitrariamente per applicazioni custom (ca.  $2^{16}$  → ca. 65000)





# Client port

---

- Il numero di porta del client è tipicamente assegnato dal sistema operativo:
  - la scelta di tale numero non è di solito rilevante.
- Quando una client socket manda un pacchetto a una server socket, il pacchetto include la specifica della porta e dell'indirizzo del client:
  - così il server è in grado di rispondere.



# Sockets protocols

---

- Usando le sockets, si può decidere che tipo di protocollo si desidera per il trasporto di pacchetti sulla rete:
  - Un protocollo **connection-oriented** (TCP) o
  - Un protocollo **connectionless** (UDP).



# Connection-oriented protocols

---

- Un protocollo connection-oriented offre l'equivalente di una conversazione telefonica.
  - Dopo aver stabilito la connessione, due applicazioni possono scambiarsi dati.
  - La connessione rimane in essere anche se nessuno parla.
- Il protocollo **garantisce** che non vengano persi dati e che questi arrivino sempre nell'ordine corretto.
- La classe Java **Socket** usa un protocollo *connection-oriented*, e “parla” **TCP**



# Connection oriented Protocols

---

- Con un connection-oriented protocol, una client socket stabilisce, alla sua creazione, una connessione con una server socket. Stabilita la connessione, un protocollo connection-oriented assicura la consegna affidabile dei dati, ovvero:
  - per ogni pacchetto spedito, il pacchetto viene consegnato.
  - ad ogni spedizione, la socket si aspetta di ricevere un acknowledgement che il pacchetto è stato ricevuto con successo.
- Se la socket non riceve l' acknowledgement entro un tempo prestabilito, la socket ri-invia il pacchetto. La socket continua a provare finché la trasmissione ha successo, o finché decide che la consegna è impossibile.
  - Ogni pacchetto IP ha un identificatore; il ricevente può quindi riordinare i pacchetti e/o rifiutare duplicati

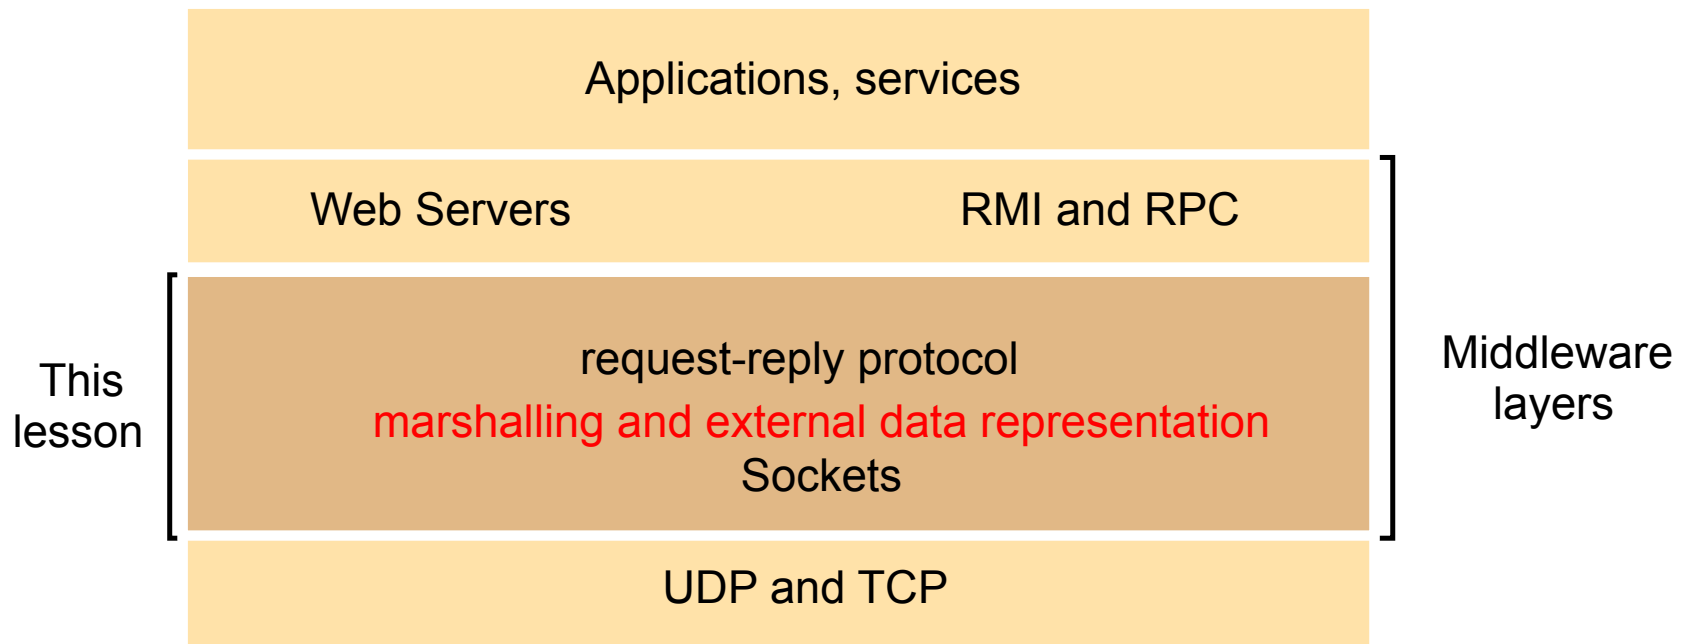


# Connectionless Protocols

---

- Un protocollo connectionless somiglia al servizio postale.
- Le applicazioni possono inviarsi brevi messaggi, ma non viene tenuta aperta una connessione tra un messaggio e l'altro.
- Il protocollo **NON garantisce** che non vadano persi dati ne' che questi arrivino nell'ordine corretto.
- La classe Java **DatagramSocket** usa un protocollo *connectionless*, e “parla” **UDP**

# Middleware layers





# Some Problems

---

- How to transfer data in different hosts
  - external data representation
  - marshalling data
  - `www` → `html`
- Define a shared communication protocol
  - request-reply protocol
  - `www` → `http`



# External Data Representation

---

- Information stores in running programs → data structures/objects
- Information in messages → sequences of bytes
- Main Methods:
  - values are converted to an a-priori agreed format and converted back on receipt (**marshalling and unmarshalling of data**)
  - values are transmitted in the sender's format, together with information of the format





# Approaches

---

- String / ASCII / UNICODE
- COM/DCOM/COM+
  - Microsoft
  - Uniform Data Transfer
- CORBA: Common Object Request Broker Architecture
  - OMG specifications (consortium)
  - Multi languages
- Java's object serialization
  - SUN
  - only Java
- Web Services
  - XML-based Open standards: XML-RPC, RESTful Web Services
  - SOAP: Simple Object Access Protocol
  - Multi languages



# CORBA Common Data Representation

- Primitive types:
  - 15 primitive types [short, long, unsigned long(4 bytes), float, double, char, boolean....]
- Constructed types
  - note that it does not deal with objects

<i>Type</i>	<i>Representation</i>
<i>sequence</i>	length (unsigned long) followed by elements in order
<i>string</i>	length (unsigned long) followed by characters in order (can also have wide characters)
<i>array</i>	array elements in order (no length specified because it is fixed)
<i>struct</i>	in the order of declaration of the components
<i>enumerated</i>	unsigned long (the values are specified by the order declared)
<i>union</i>	.type tag followed by the selected member



# Java serialized form

<i>Serialized values</i>				<i>Explanation</i>
Person	8-byte version number		h0	<i>class name, version number</i>
3	int year	java.lang.String name:	java.lang.String place:	<i>number, type and name of instance variables</i>
1934	5 Smith	6 London	h1	<i>values of instance variables</i>

- h0 and h1 are handles (references for internal objects: i.e. String,..)
- deep copy of both public and private variable
- use of transient keyword to avoid serialization of specific fields
- **reflection**: ability to enquire class properties and create classes from names + default constructor



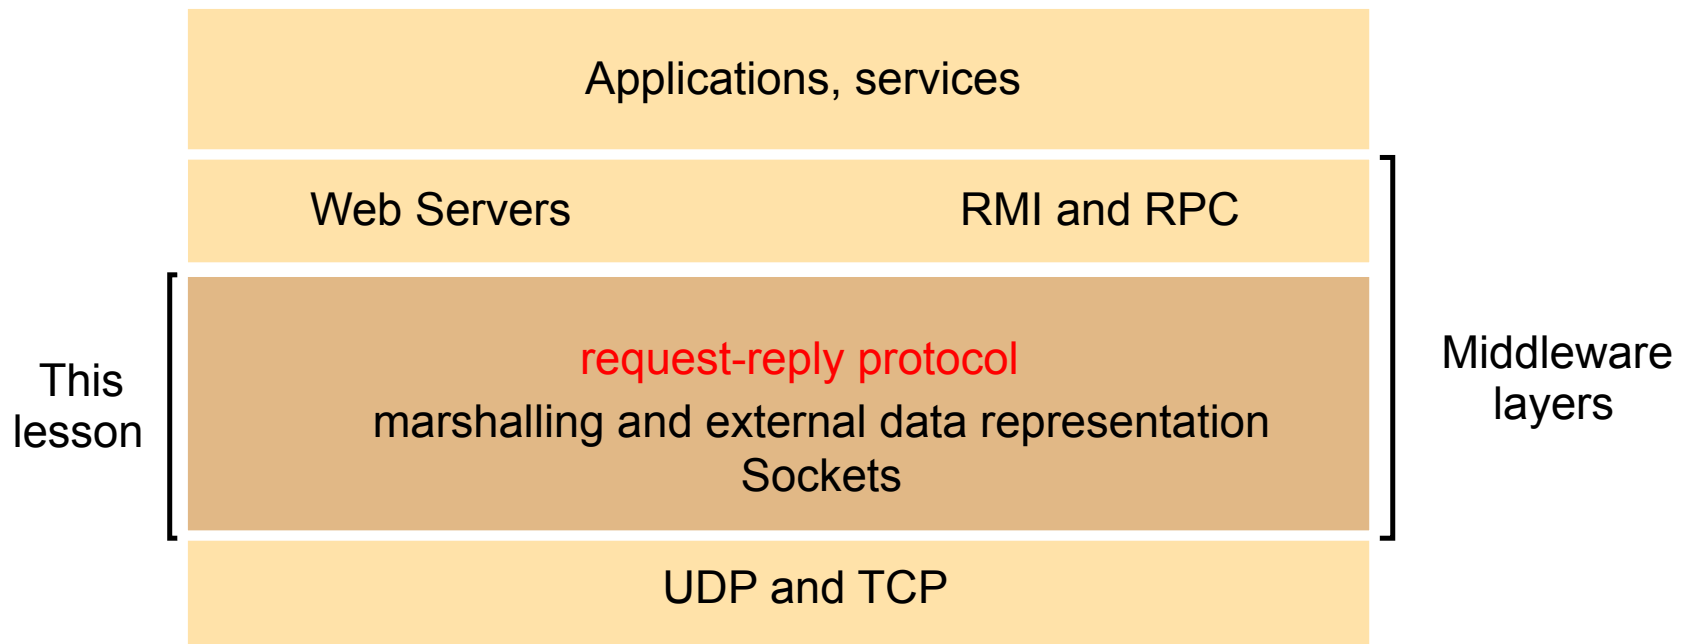
# World Wide Web

## ■ HTML - XML

GET /search  
Host: [www](#)  
User-Agent: Mozilla/5.0 (Macintosh; U; Intel Mac OS X 10\_5\_8; en-us; rv:1.9.2.15) Gecko/20110309 Firefox/3.6.15  
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,\*/\*;q=0.8  
Accept-Language: en-us  
Accept-Encoding: gzip, deflate  
Accept-Charset: utf-8, iso-8859-1, latin1  
**HTTP/1.1 200 OK**  
**Date:** Fri, 17 Sep 2009 07:59:01 GMT  
**Server:** Apache/2.0.50 (Unix) mod\_perl/1.99\_10 Perl/v5.8.4  
**mod\_ssl/2.0.50 OpenSSL/0.9.7d DAV/2 PHP/4.3.8 mod\_bigwig/2.1-3**  
**Last-Modified:** Tue, 24 Feb 2009 08:32:26 GMT  
**ETag:** "ec002-afa-fd67ba80"  
**Accept-Ranges:** bytes  
**Content-Length:** 2810  
**Content-Type:** text/html

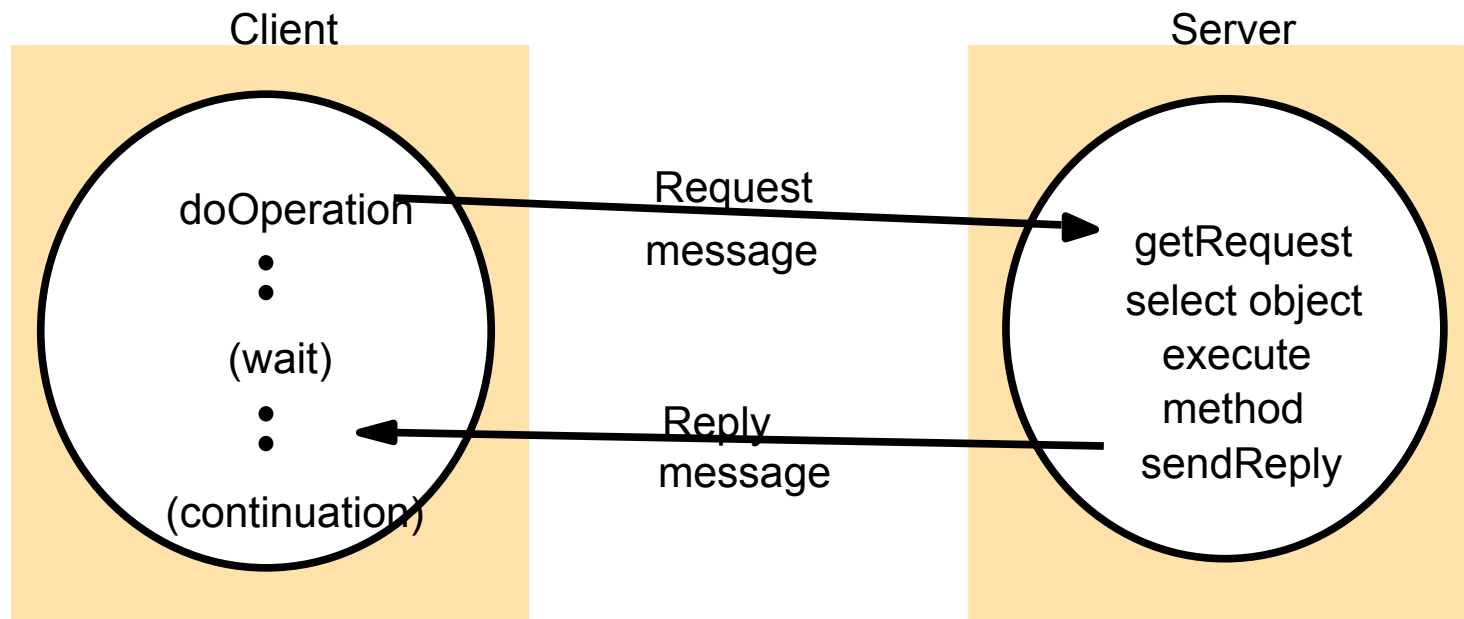
```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01 Transitional//EN">
<html>
...
....
</html>
```

# Middleware layers



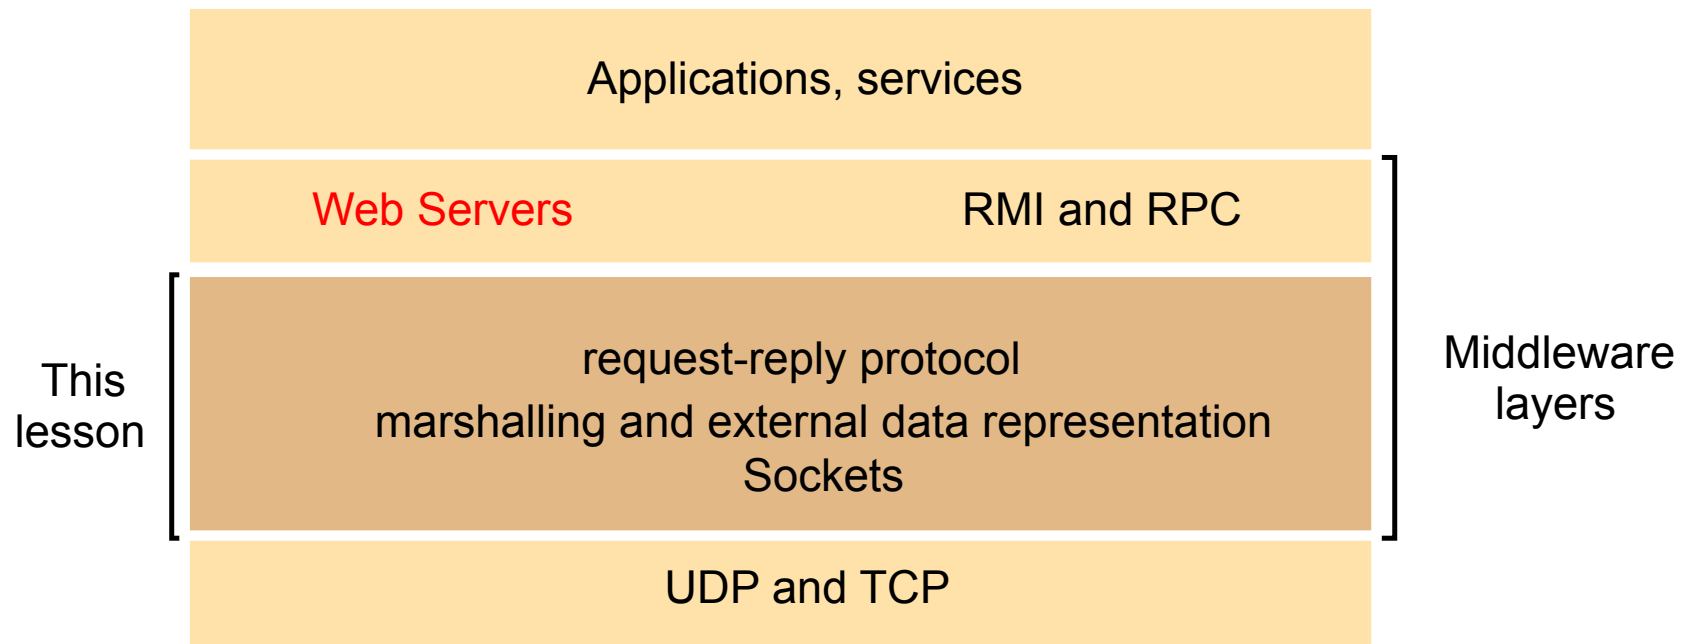
# Request-reply communication

- The majority of communication protocols (RPC/RMI/HTTP..) support such communication schema



# Middleware layers

*Maurizio Marchese – Intro Web Prog*



# Brief History of HTTP

Maurizio Marchese – Intro Web Prog


- 1990, the Http protocol was proposed by Tim Bernes Lee and was used at CERN
- Summer of 1991, the http protocol was released to the public.
- 1993: Marc Andreessen (then a grad student at NSCA) posts Mosaic on an ftp site. New features include:
  - Hyperlinks
  - Embedded images
- December 1993: Mosaic growth makes the front page of New York Times
- 1994: Marc Andreessen and colleagues leave NSCA to form Mosaic Corp. (later renamed "Netscape")





# Web Growth


Maurizio Marchese – Intro Web Prog



# Web Server Statistics

- Apache is the most popular web server today (freely available)
- Microsoft IIS is gaining ground

Market Share for Top Servers Across All Domains August 1995 - April 2009





# Versions of HTTP

---

- Early protocol is HTTP 0.9
  - read only
- More recent versions:
  - HTTP 1.0
    - read, input, delete, ...
  - HTTP 1.1
    - performance optimizations



# HTTP Overview

---

- Client (browser) sends HTTP request to server
- Request specifies affected *URL*
- Request specifies desired *operation*
- Server performs *operation* on *URL*
- Server sends response
- Request and reply headers are in pure text



# Static Content and HTML

---

- | ■ Most static web content is written in HTML
- HTML allows
  - Text formatting commands
  - Embedded objects
  - Links to other objects
- Server do not need to understand or interpret HTML; Browsers do



# HTTP Request Operations

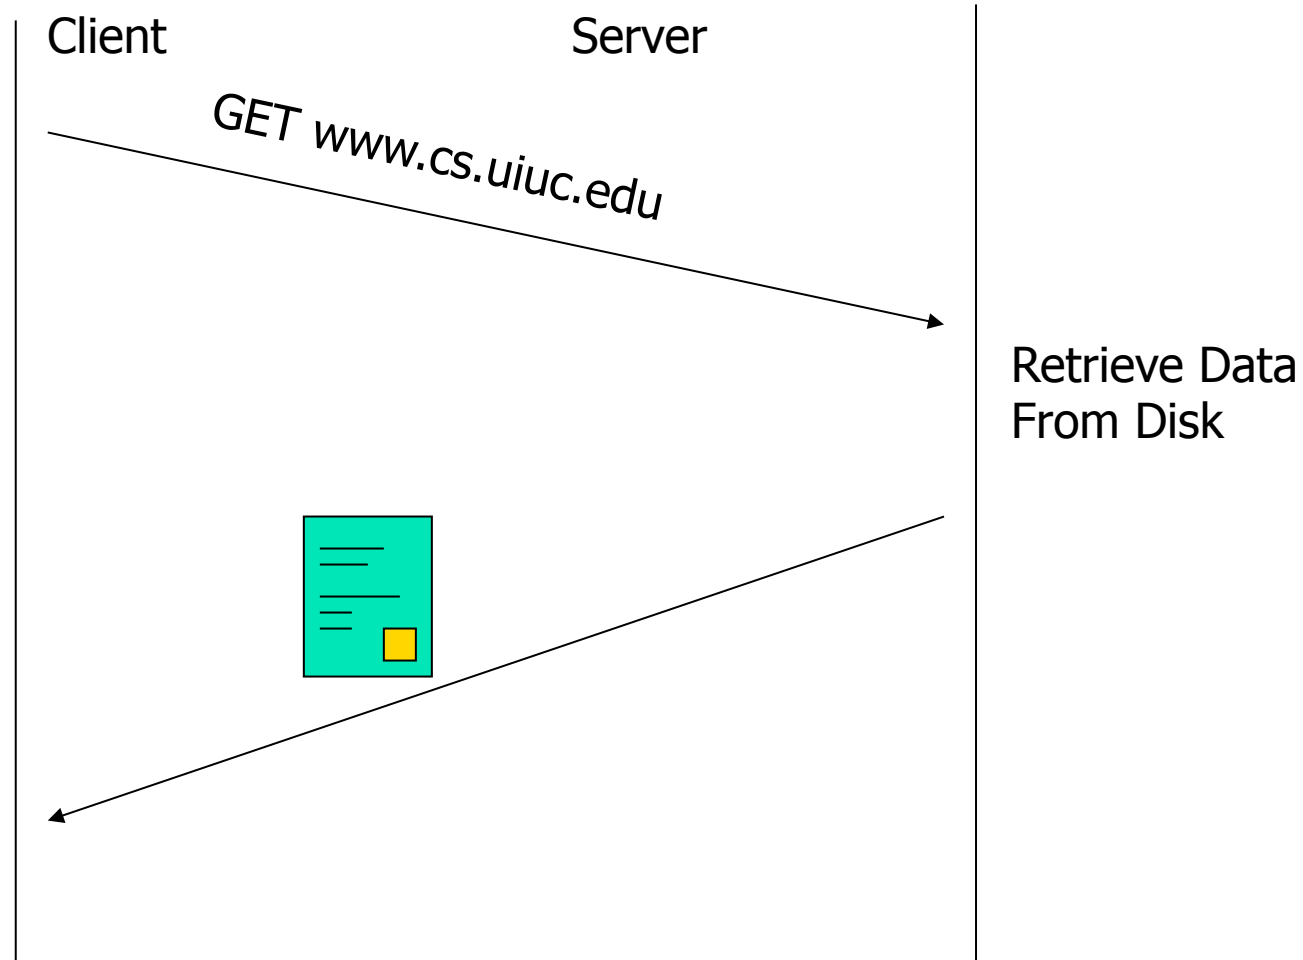
---

- GET: retrieves URL (most widely used)
- HEAD: retrieves only response header
- POST: posts data to server
- PUT: puts page on server
- DELETE: deletes page from server

more later...

# Example of an HTTP 1.0 Exchange

Maurizio Marchese – Intro Web Prog





# Fetching Multiple Objects (1.0)

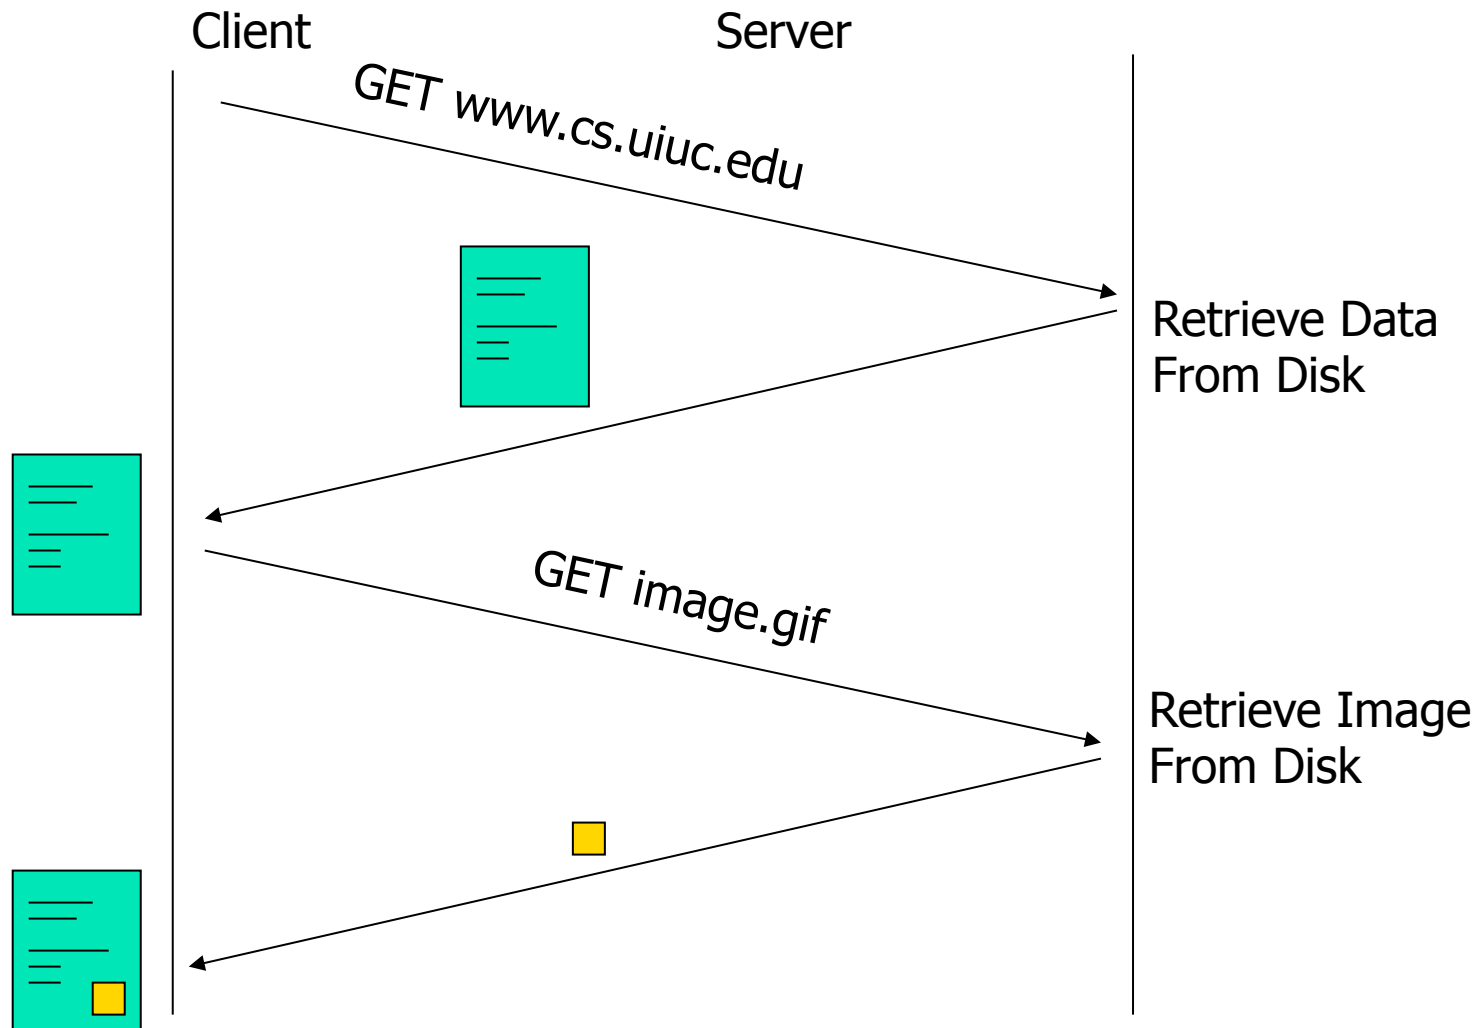
---

- Most web-pages contain embedded objects (e.g., images, backgrounds, etc)
- Browser requests HTML page
- Server sends HTML file
- Browser parses file and requests embedded objects
- Server sends requested objects



# Fetching Embedded Objects (1.0)

Maurizio Marchese – Intro Web Prog





# HTTP Reply

---

- The http response from the server to the client is sent using the same TCP connection as the request (**NB TCP connection are bidirectional**)
- It consist of an **Header** and a **Body**
  - **Header**: Status, Date, Server info, Body info..
  - **Body**: if request is successful contains the resource that has been requested

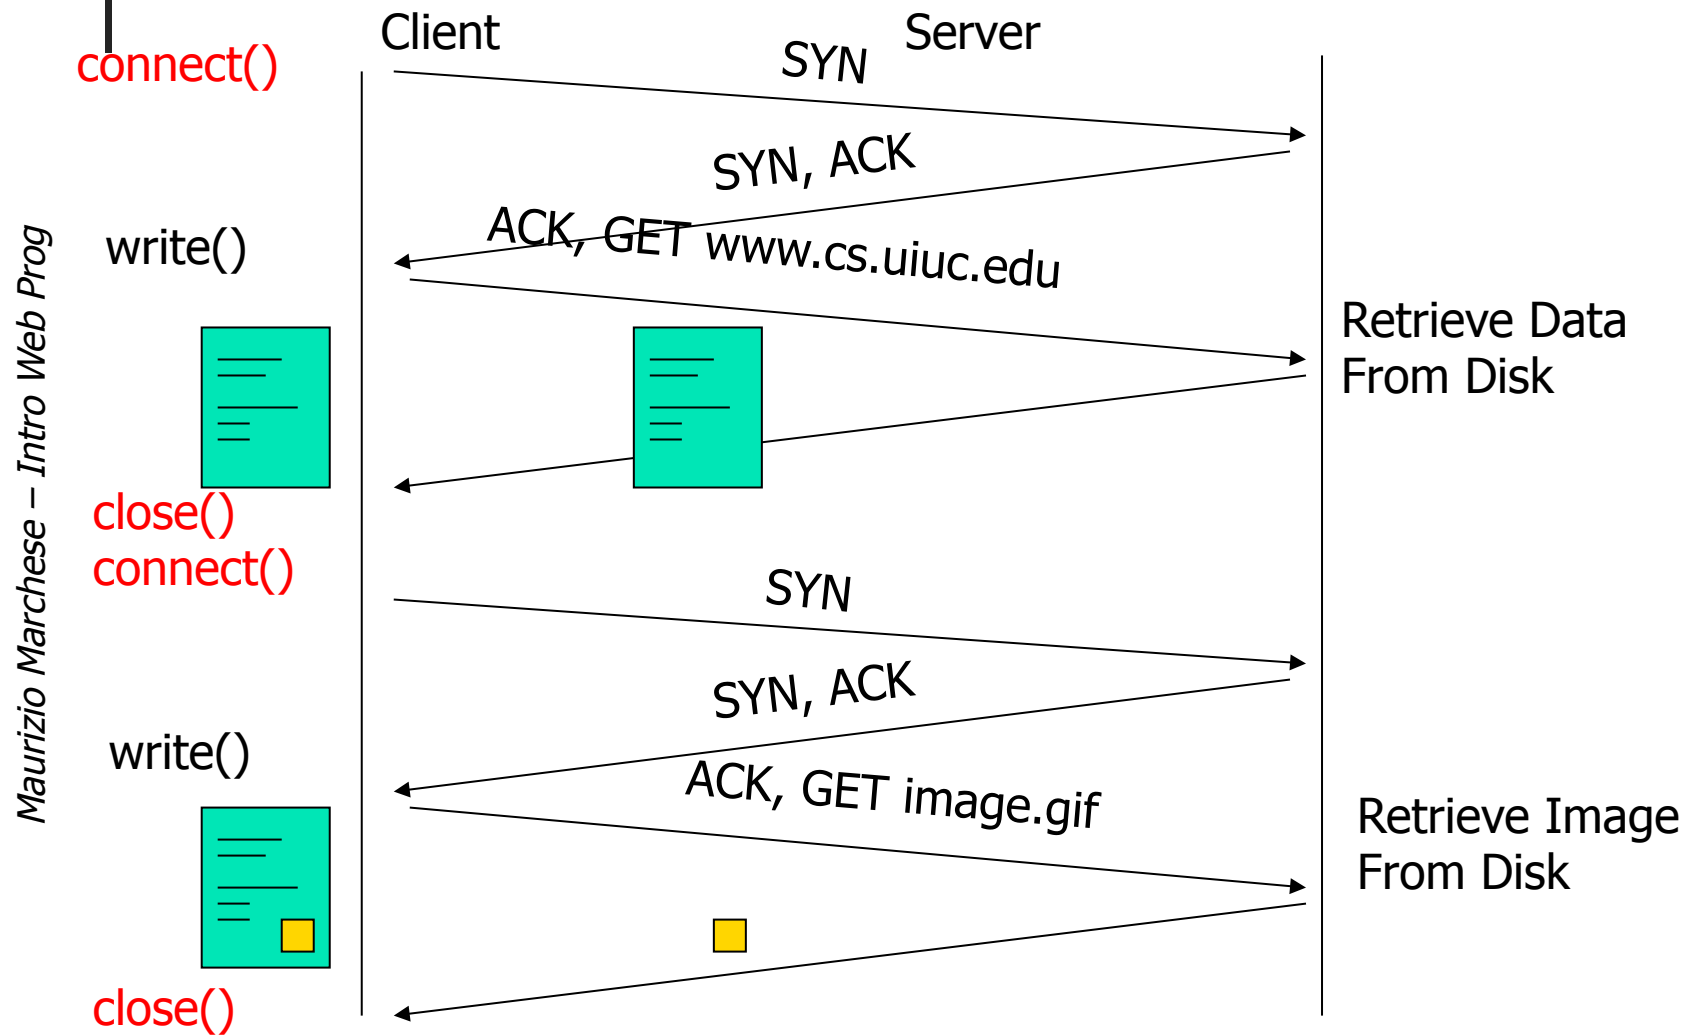


# HTTP 1.0

---

- Client opens a separate TCP connection for each requested object
- Object is served and connection is closed
- Advantages
  - maximum concurrency
- Limitations
  - TCP connection setup/tear-down overhead
  - TCP slow start overhead

# HTTP 1.0





# HTTP 1.1

---


- To avoid a connection per object model, HTTP 1.1 supports *request persistent connections*
- Client opens TCP connection to server
- All requests use same connection
- When the whole page (including sub-objects) is loaded the TCP connection is closed
- Problems
  - Less concurrency
  - Server does not know when to close idle connections

# HTTP 1.1


Maurizio Marchese – Intro Web Prog

connect()

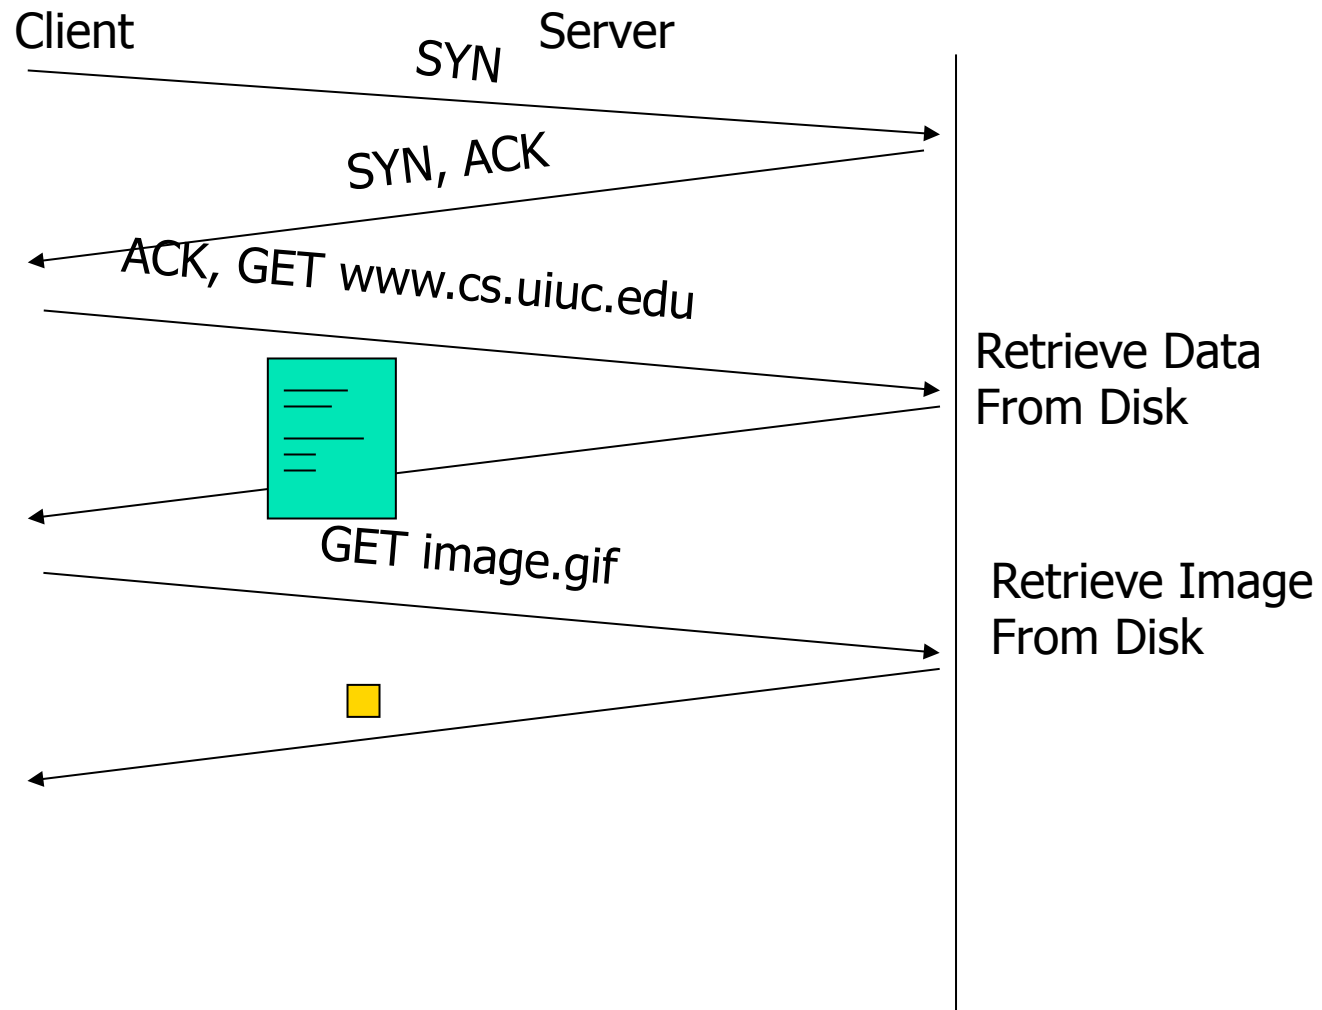
write()



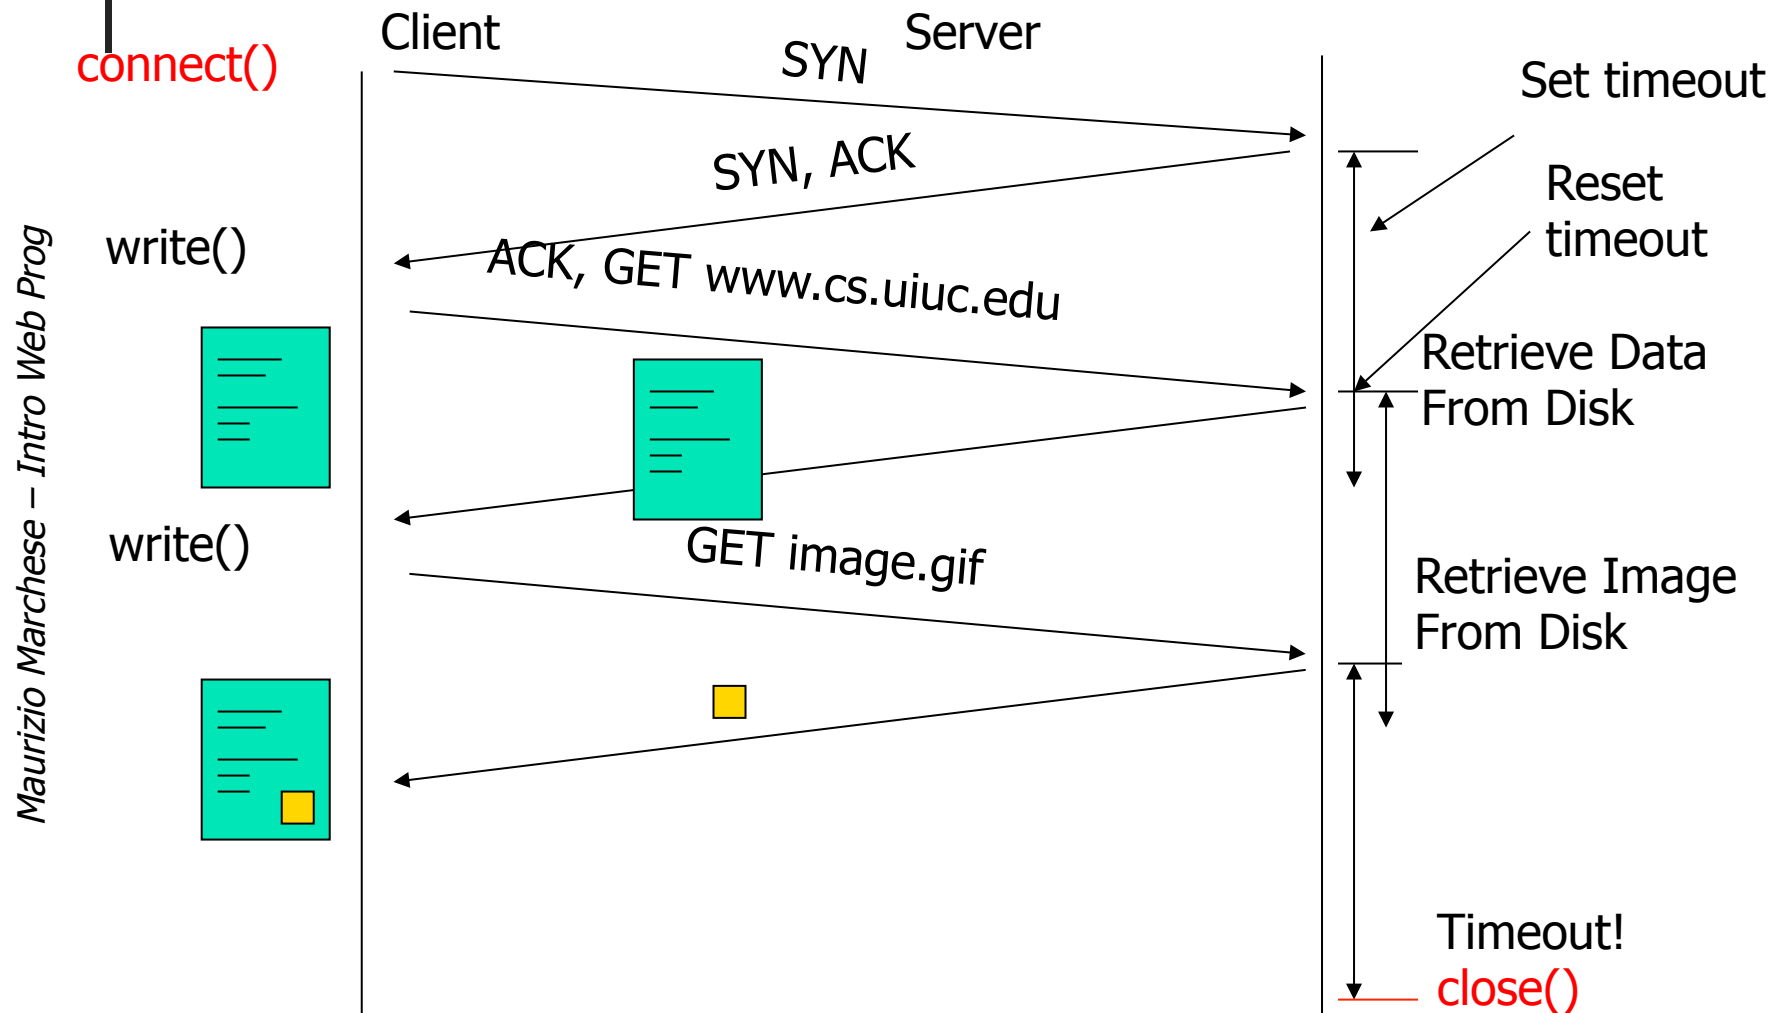
write()



close()



# Server Side Close()





# HTTP Operations

---

- GET: retrieves URL (most widely used)
- HEAD: retrieves only response header
- POST: posts data to server
- PUT: puts page on server
- DELETE: deletes page from server



# HTTP request-replay message

## ■ request message

<i>method</i>	<i>URI or pathname</i>	<i>HTTP version</i>	<i>headers</i>	<i>message body</i>
GET	//www.google.com	HTTP/ 1.1	.....	.....

## ■ replay message

<i>HTTP version</i>	<i>status code</i>	<i>reason</i>	<i>headers</i>	<i>message body</i>
HTTP/1.1	200	OK	.....	resource data



# URI = URL+URN

---

- URI *Uniform Resource Identifier* is a standard mechanism to identify electronic resources
- URI = *Uniform Resource Locator + Uniform Resource Name*
- Syntax:  
`<scheme>` : <scheme dependent info>
- Example:  
`http` : `//www.isoc.org/internet/history/`



# HTTP Request

---

- HTTP uses Uniform Resources Identifier (URI)

- abstract URI: <http://host:port/path?query>
- real URI: <http://www.google.com/search?q=An+Introduction+to+Web+Technologies>

- The HTTP request is sent using TCP.

- in the above example, the browser will create a message like the following

GET /search?q=An+Introduction+to+Web+Technologies HTTP/1.1

Host: [www.google.com](http://www.google.com)

User-Agent: Mozilla/5.0 (X11; U; Linux i686; en-US; rv1.7.2)

Accept: text/xml, application/xml, application/xhtml+xml, text/html;q=0.9, text/plain;q=0.8, image/png, \*/\*;q=0.5

Accept-Language: it, en-us;q=0.8,en;q=0.5,sw;q=0.3

Accept-Encoding: gzip,deflate

Accept-Charset: ISO-8859-1, utf-8; q=0.7,\*; q=0.7

...



# HTTP GET anatomy (1)

---

**GET /search?q=An+Introduction+to+Web+Technologies HTTP/1.1**

- ask the server to the resource "search?q=An+Introduction+to+Web+Technologies" using version 1.1. of the HTTP protocol
- It uses a GET because the creator of the web page has programmed so in the code
- all other lines are header lines each having the form field:value

**Host: www.google.com**

- field: Host
- value: domain name of the server

**User-Agent: Mozilla/5.0 (X11; U; Linux i686; en-US; rv1.7.2)**

- value: information about the user agent (browser or search robot) to tailor the response



# HTTP GET anatomy (2)

---

**Accept: text/xml, application/xml, application/xhtml+xml, image/png, text/html;q=0.9, text/plain;q=0.8, \*/\*; q=0.5**

→ value: specify media types that are acceptable as a response (MIME - Multipurpose Internet Mail Extensions - types)

The q parameters indicate a relative quality factor (0: not acceptable; 1:default)

**Accept-Language: it, en-us;q=0.8,en; q=0.5,sw; q=0.3**

→ value: acceptable/preferred natural languages

**Accept-Encoding: gzip,deflate**

→ value: acceptable/preferred content codings

**Accept-Charset: ISO-8859-1, utf-8; q=0.7,\*; q=0.7**

→ value: acceptable/preferred character encoding



# HTTP Response anatomy

---

**HTTP/1.1 200 OK**

**Date:** Fri, 17 Sep 2009 07:59:01 GMT

**Server:** Apache/2.0.50 (Unix) mod\_perl/1.99\_10 Perl/v5.8.4  
mod\_ssl/2.0.50 OpenSSL/0.9.7d DAV/2 PHP/4.3.8 mod\_bigwig/2.1-3

**Last-Modified:** Tue, 24 Feb 2009 08:32:26 GMT

**ETag:** "ec002-afa-fd67ba80"

**Accept-Ranges:** bytes

**Content-Length:** 2810

**Content-Type:** text/html

HEADER

<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01 Transitional//EN">

<html>

...

....

</html>

BODY



# Status Codes

---

- 200 OK
- 301 Moved Permanently
- 400 Bad Request
- 401 Unauthorized
- 403 Forbidden
- 404 Not Found
- 500 Internal Server Error
- 503 Service Unavailable



# Other HTTP methods: POST

---

- The HTTP-GET protocol creates a query string of the name-and-value pairs and then appends the query string to the URL of the script on the server that handles the request.
  - Therefore, you can “read” easily the request.
- The HTTP-POST protocol passes the name-and-value pairs in the body of the HTTP request message. It is useful:
  - to provide data to data-handling process (servlet,cgi,php,...)
  - to interact with databases
  - to improve (not much) privacy





# Limitations of HTTP

---

- **Stateless**, no built-in support for tracking clients (session management)
  - Need for special programming techniques for session management: Cookies, Sessions, URL rewriting, Hidden Form Fields.
- No built-in **security** mechanisms
  - Need for special security protocols



# Security

---

- **SSL**: *Secure Sockets Layer*
- **TLS**: *Transport Layer Security* (newer version)
- Layer between HTTP and TCP, accessed by **HTTPS://**...
- Based on public-key cryptography
  - private key + public key
  - certificate (usually for server authentication only)
- The main idea of HTTPS is to create a secure channel over an insecure network.
  - This ensures reasonable protection from eavesdroppers and man-in-the-middle attacks, provided that adequate cipher suites are used and that the server certificate is verified and trusted.
- The trust inherent in HTTPS is based on major certificate authorities that come pre-installed in browser software
  - this is equivalent to saying "I trust certificate authority (e.g. VeriSign/Microsoft/etc.) to tell me whom I should trust"



# Dynamic Content

---

- Web pages can be created as requests arrive
- Advantages
  - Personalization (e.g., my.yahoo.com),
  - interaction with client input
  - interaction with back-end applications
- Disadvantages
  - Performance penalty
- Generating dynamic content (CGI, ASP, PHP, ColdFusion, JavaScript, Flash, ...)



# CGI Scripts

---

- CGI scripts are URLs with a .cgi extension
- The script is a program (e.g., C, JAVA, ...)
- When the URL is requested, server invokes the named script, passing to it client info
- Script outputs HTML page to standard output (redirected to server)
- Server sends page to client

# CGI Execution

*Maurizio Marchese – Intro Web Prog*

