
DDL / SQL

**Data Description Language /
Data Definition Language
Structured Query Language**

Wikipedia: DDL

- A data definition language or data description language (DDL) is a syntax similar to a computer programming language for defining data structures, especially database schemas.
- The data definition language concept and name was first introduced in relation to the Codasyl database model, where the schema of the database was written in a language syntax describing the records, fields, and sets of the user data model.
- Later it was used to refer to **a subset of Structured Query Language (SQL) for creating tables and constraints.**
- These information tables were specified as SQL/Schemata in SQL:2003.
- The term DDL is also used in a generic sense to refer to any formal language for describing data or information structures.

Tools

- Open Source DBMS

- MySQL

- ◆ for windows →

- ✓ <http://www.mysql.com/why-mysql/windows/>

- ◆ For Unix →

- ✓ <http://dev.mysql.com/doc/mysql-linuxunix-excerpt/5.1/en/index.html>

- ◆ For Mac →

- ✓ <http://www.mamp.info/en/index.html>

- Java DB

- ◆ Oracle's supported distribution of the Apache Derby open source database

- ◆ Integrated in Netbeans

- ◆ Tutorial:

- ◆ <http://www.w3schools.com/sql/>

Creating a Database Table

Syntax: **CREATE TABLE** <name> (<list of elements>)

An element is an attribute definition of the form

attributeName Type

Type can be: Integer, CHAR(XX), VARCHAR(XX), Date, Real

Do not forget the comma between attribute definitions

Important the difference between CHAR and VARCHAR

KEY (<attribute list>)

FOREIGN KEY (<attribute list>) REFERENCES <Table>

Example: *Studio*(name, address, budget, president), Person(name, tel)

```
CREATE TABLE Studio (  
    name CHAR(20),  
    address VARCHAR(20),  
    budget REAL,  
    president CHAR(20),  
    PRIMARY KEY (name),  
    FOREIGN KEY (president) REFERENCES Person(name) )
```

```
CREATE TABLE Person (  
    name CHAR(20),  
    tel INTEGER,  
    PRIMARY KEY (name) )
```

Deleting a Database Table

Syntax: **DROP TABLE** <name>

Example:
DROP TABLE *Studio*

Types

1. **INT** or **INTEGER**.
2. **BOOLEAN**
3. **REAL** or **FLOAT**.
4. **CHAR**(n) = fixed length character string, padded with “pad characters.”
5. **VARCHAR**(n) = variable-length strings up to n characters.
6. **DATE & TIME**

Declaring Keys

Use **PRIMARY KEY** or **UNIQUE**.

- only one **PRIMARY KEY**
 - ◆ The primary key of a relational table uniquely identifies each record in the table. It can either be a normal attribute that is guaranteed to be unique (such as Social Security Number in a table with no more than one record per person) or it can be generated by the DBMS
- many **UNIQUEs** allowed.
- SQL permits implementations to create an *index* (data structure to speed access given a key value) in response to **PRIMARY KEY** only.
- SQL does not allow nulls in primary key, but allows them in “unique” columns (which may have two or more nulls, but not repeated non-null values).
- Each table should have a primary key, and each table can have only **ONE** primary key.

Example

Find the differences...

```
CREATE TABLE Studio-version1 (  
    name CHAR(20),  
    address VARCHAR(20) UNIQUE,  
    president REAL UNIQUE,  
    PRIMARY KEY(name))
```

```
CREATE TABLE Studio-version2 (  
    name CHAR(20),  
    address VARCHAR(20),  
    president REAL,  
    UNIQUE (address, president),  
    PRIMARY KEY(name))
```


Other Properties You Can Give to Attributes

1. **NOT NULL** = every tuple must have a real value for this attribute.
2. **DEFAULT** value = a value to use whenever no other value of this attribute is known.

Example:

```
CREATE TABLE Studio (  
    name CHAR(30) PRIMARY KEY,  
    addr CHAR(50) DEFAULT '123 Sesame St',  
    president_id CHAR(50) NOT NULL)
```

Example

Consider executing the following statement

```
INSERT INTO Studio(name, president_id )  
VALUES('Fox', 90943)
```

The result is the following tuple:

<i>name</i>	<i>addr</i>	<i>president_id</i>
Fox	123 Sesame St.	90943

Primary key is by default not NULL.

This insertion is legal. It is OK to list a subset of the attributes and values for only this subset. But if we had forgot to specify a value for *president_id* then the insertion could not be made.

Changing Columns

Add an attribute to relation R with

ALTER TABLE R ADD <column declaration>

Example:

CREATE TABLE *MovieStar* (

name **CHAR**(30),

address **VARCHAR**(255),

gender **CHAR**(1),

birthdate **DATE**,

PRIMARY KEY(*name*))

ALTER TABLE *MovieStar* **ADD** *phone* **CHAR**(16) **DEFAULT** 'unlisted'

Columns may also be dropped...

Example"

ALTER TABLE *MovieStars* **DROP** *birthdate*

Database Modifications

CRUD operations: Create (insert), Read (select), Update(update), Delete (delete)

e.g. to insert a tuple we use the statement

INSERT INTO <relation> **VALUES** (<list of values>)

- In the above statement the values as listed in the same order with which the attributes were declared
- If we want to ignore this order then we list the attributes as arguments of the relation

Example:

Consider *StarsIn*(movieTitle, movieYear, starName)

Insert the fact that Sydney Green stars in The Maltese Falcon

INSERT INTO *StarsIn*(*movieTitle*, *starName*, *movieYear*)
VALUES('The Maltese Falcon', 'Sydney Green', 1942);

Insertion of the Result of a Query

Syntax: **INSERT INTO** <relation> (<subquery>).

(not so important for the Intro Web Prog. course ...)

Example:

Consider

Studio(*name*, *address*, *presC#*)

Movie(*title*, *year*, *length*, *inColor*, *studio-Name*, *prodC#*)

Add to relation *Studio* all the studios that are mentioned in relation *Movie*

```
INSERT INTO Studio(name)  
  SELECT DISTINCT studio-Name  
  FROM Movie  
  WHERE studio-Name NOT IN (SELECT name  
                                FROM Studio)
```

Deletion

Syntax: **DELETE FROM** <relation> **WHERE** <condition>

Semantic: Deletes all tuples satisfying the condition from
the named relation

Example:

Consider *StarsIn*(movieTitle, movieYear, starName)

Sydney Green was not a star in The Maltese Falcon...

```
DELETE FROM StarsIn
WHERE movieTitle = 'The Maltese Falcon' AND
      movieYear = 1942 AND starName = 'Sydney Green'
```

As another example, to make the *StarsIn* relation empty execute

```
DELETE FROM StarsIn
```

Updates

Syntax: **UPDATE** <relation>

SET <new-value assignments>

WHERE <condition>

Example:

Consider

Customers(CustID, CustName, ContactName, Address, City, PostalCose, Country)

```
UPDATE Customers  
SET ContactName = 'Pippo ', City = 'Trento'  
WHERE CustName = 'Pluto'
```



SQL



SQL Queries

Principal form:

SELECT desired attributes

FROM tuple variables — range over relations

WHERE condition about tuple variables

Running example relation schema:

Movie(title, year, length, inColor, studio-Name, producerC#)

StarsIn(movieTitle, movieYear, starName)

MovieStar(name, address, gender, birthdate)

MovieExec(name, address, cert#, netWorth)

Studio(name, address, presC#)

Example

Consider relation:

Movie(title, year, length, inColor, studio-Name, producerC#)

Find all the movies produced by Disney Studios in 1990...

```
SELECT *  
FROM Movie  
WHERE studio-Name='Disney' AND year=1990
```

Find all the movies made by Fox that are at least 100 min. long

```
SELECT *  
FROM Movie  
WHERE studio-Name='Fox' AND length >= 100
```

Star as List of All Attributes

	<u>title</u>	year	length	in-Color	studio-Name	produceC#
Movie:	Star Wars	1977	124	true	Fox	12345
	Mighty Ducks	1991	104	true	Disney	67890
	Wayne's World	1992	95	true	Paramaount	99999
	Spider-Man	2002	121	true	Columbia	12345
	Episode I	1999	133	true	Fox	45634
	Episode II	2002	142	true	Fox	23456

SELECT *

FROM *Movie*

WHERE *studio-Name*='Disney' **AND** *year*>=1990

<u>title</u>	year	length	in-Color	studio-Name	produceC#
Mighty Ducks	1991	104	true	Disney	67890

Projection in SQL

	<i>title</i>	<i>year</i>	<i>length</i>	<i>in-Color</i>	<i>studio-Name</i>	<i>produceC#</i>
<i>Movie:</i>	Star Wars	1977	124	true	Fox	12345
	Mighty Ducks	1991	104	true	Disney	67890
	Wayne's World	1992	95	true	Paramaount	99999
	Spider-Man	2002	121	true	Columbia	12345
	Episode I	1999	133	true	Fox	45634
	Episode II	2002	142	true	Fox	23456

SELECT *title, length*

FROM *Movie*

WHERE *studio-Name*='Disney' **AND** *year*>=1990

<i>title</i>	<i>length</i>
Mighty Ducks	104

Expressions as Values in Columns

<i>Movie:</i>	<u><i>title</i></u>	<i>year</i>	<i>length</i>	<i>in-Color</i>	<i>studio-Name</i>	<i>produceC#</i>
	Star Wars	1977	124	true	Fox	12345
	Mighty Ducks	1991	104	true	Disney	67890
	Wayne's World	1992	95	true	Paramaount	99999
	Spider-Man	2002	121	true	Columbia	12345
	Episode I	1999	133	true	Fox	45634
	Episode II	2002	142	true	Fox	23456

SELECT *title AS name, length * 0.016667 AS h-duration*
FROM *Movie*
WHERE *studio-Name='Disney' AND year>=1990*

<u><i>name</i></u>	<u><i>h-duration</i></u>
Mighty Ducks	1

Patterns

- % stands for any string.
- _ stands for any one character.
- “Attribute **LIKE** pattern” is a condition that is true if the string value of the attribute matches the pattern. Also **NOT LIKE** for negation.

Example:

Given relation *Movie*(*title*, *year*, *length*, *inColor*, *studio-Name*, *producerC#*)

find movies with title “Star *something*”, where *something* has 4 letters...

```
SELECT title  
FROM Movie  
WHERE title LIKE 'Star _ _ _ _'
```

Note: patterns must be quoted, like strings.

Nulls

Used in place of a value in a tuple's component.

- Interpretation is not exactly “missing value.”
- There could be many reasons why no value is present, e.g., “value inappropriate.”

Operations with expressions that evaluate to **NULL**

- The result is always **NULL**

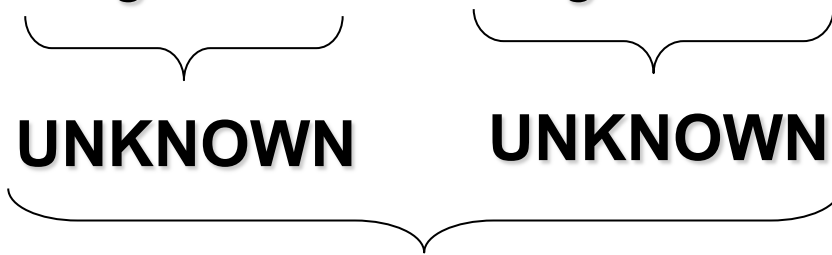
Comparing expressions that evaluate to **NULL** with values

- 3rd truth value **UNKNOWN**.
- A query only produces tuples if the **WHERE**-condition evaluates to **TRUE** (**UNKNOWN** is not sufficient).

Example

<u>title</u>	year	length	in-Color	studio-Name	produceC#
Mighty Ducks	1991	NULL	true	Disney	67890

SELECT *title, year*
FROM *Movie*
WHERE *length < 90 OR length >= 90*



UNKNOWN

Mighty Ducks is not *selected*, even though
the **WHERE** condition is a tautology.

Multi-relation Queries

- List of relations in **FROM** clause.
- Relation-dot-attribute disambiguates attributes from several relations

Example:

Consider

Movie(title, year, length, inColor, studio-Name, producerC#)


MovieExec(name, address, cert#, netWorth)

I wonder who is the producer of Star Wars...

```
SELECT name  
FROM Movie, MovieExec  
WHERE title = 'Star Wars' AND producerC# = cert#
```

Example

```
SELECT name  
FROM Movie, MovieExec  
WHERE title = 'Star Wars' AND producerC# = cert#
```



Explicit Tuple Variables

Sometimes we need to refer to two or more copies of a relation.

To do that we use *tuple variables* as aliases of the relations.

Example:

Consider *MovieStar*(*name*, *address*, *gender*, *birthdate*)

Now, find two stars that share the same address....

```
SELECT Star1.name, Star2.name  
FROM MovieStar Star1, MovieStar Star2  
WHERE Star1.address = Star2.address AND  
       Star1.name < Star2.name
```

Note that *Star*₁.*name* < *Star*₂.*name* is needed to avoid producing (Carrie, Carrie) and to avoid producing a pair in both orders.

Union/Intersection/Difference

Consider

MovieStar(name, address, gender, birthdate)

MovieExec(name, address, cert#, netWorth)

Find all the names and addresses of people that are either movie stars **or** movie executives

```
(SELECT name, address
FROM MovieStar)
UNION
(SELECT name, address
FROM MovieExec)
```

Union/Intersection/Difference

Consider

MovieStar(name, address, gender, birthdate)

MovieExec(name, address, cert#, netWorth)

Find all the female movie stars who are also executives **and** have a net worth over \$10,000,000

```
(SELECT name, address
FROM MovieStar
WHERE genter = 'F')
INTERSECT
(SELECT name, address
FROM MovieExec
WHERE netWorth > 10000000)
```

Union/Intersection/Difference

Consider

MovieStar(name, address, gender, birthdate)

MovieExec(name, address, cert#, netWorth)

Find all the movie stars who **are not** executives

```
(SELECT name, address
FROM MovieStar)
EXCEPT
(SELECT name, address
FROM MovieExec)
```

Forcing Set/Bag Semantics

- Default for select-from-where is bag semantics
i.e., duplicate rows are retained.
 - ◆ Why? Saves time of not comparing tuples as we generate them.
- Default for union, intersection, and difference is set semantics
i.e. a set in mathematics is a collection of well defined and distinct object
 - ◆ we need to sort anyway when we take intersection or difference.
(Union seems to be thrown in for good measure!)
- Force set semantics with **DISTINCT** after **SELECT**.
 - ◆ But make sure the extra time is worth it.
- Force bag semantics in Union/Intersection/Difference using **ALL**

Example

Consider

MovieStar(name, address, gender, birthdate)

MovieExec(name, address, cert#, netWorth)

Find all the names of people that are either movie stars or movie executives

```
(SELECT name, address
FROM MovieStar)
UNION ALL
(SELECT name, address
FROM MovieExec)
```

Find all the unique names of people that are stars

```
SELECT DISTINCT name
FROM MovieStar
```